

# Machine Learning

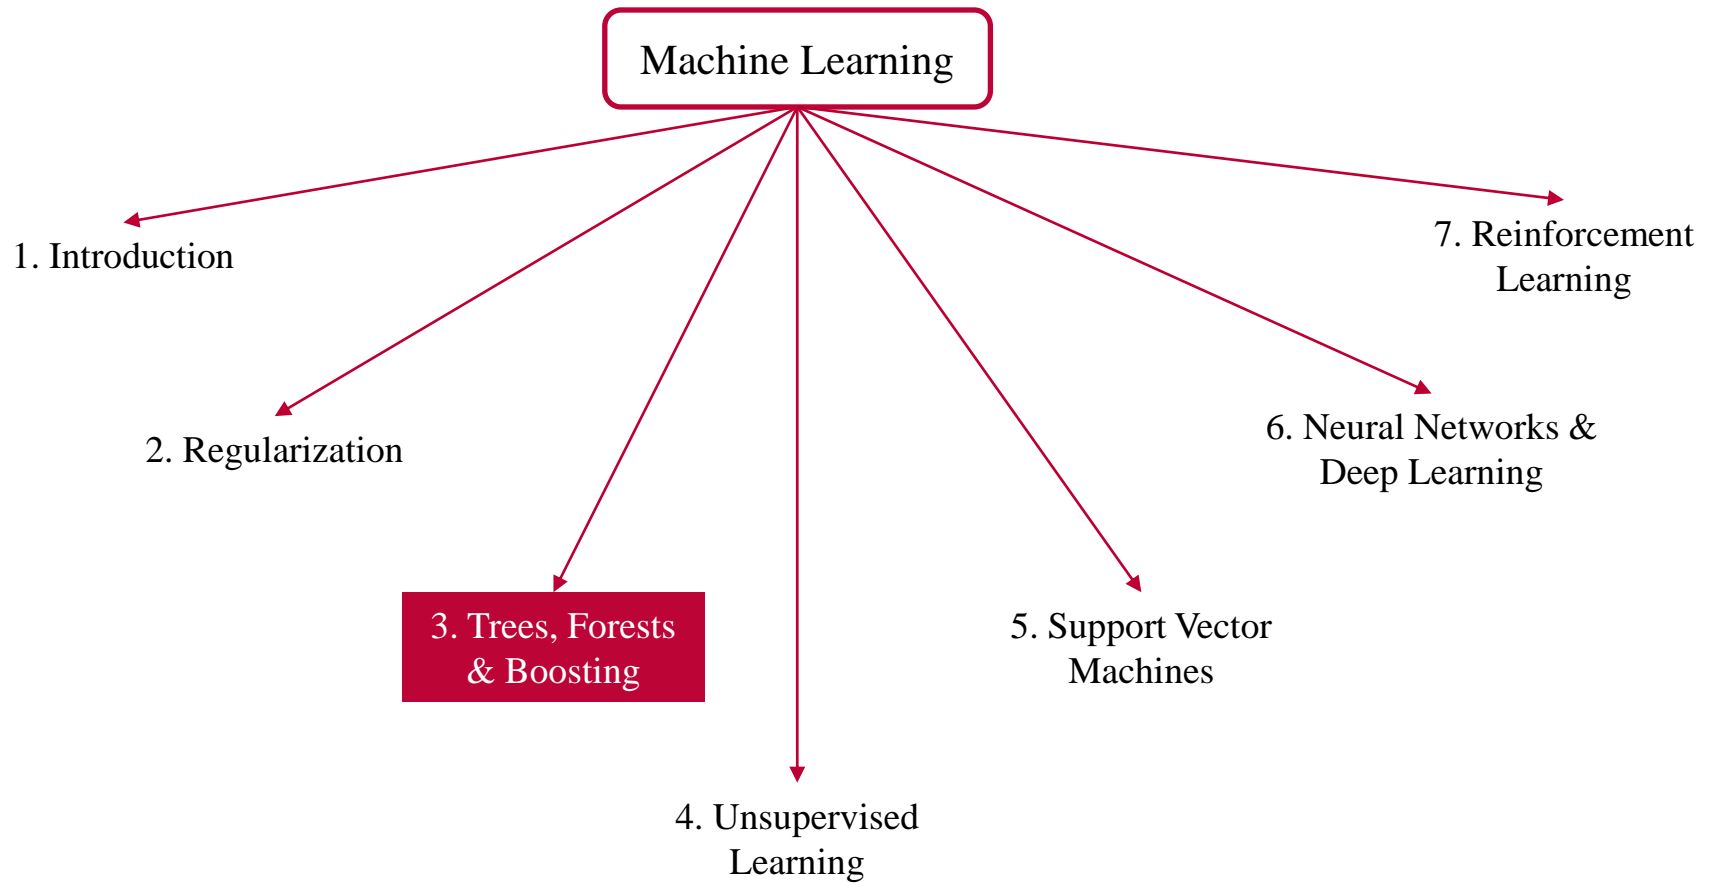
FEM31002

## Trees, Forests and Boosting

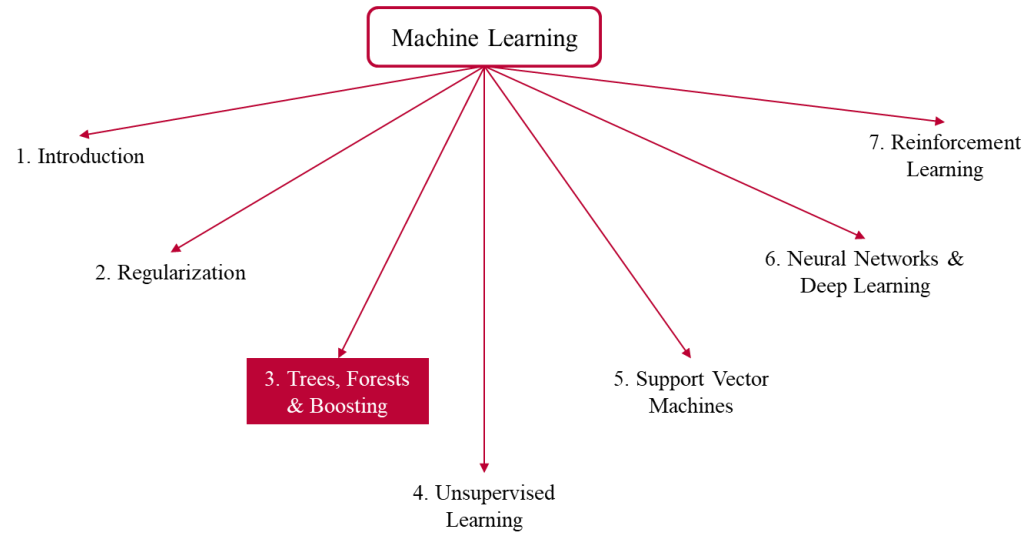
**Ilker Birbil**

[birbil@ese.eur.nl](mailto:birbil@ese.eur.nl)

# Outline



# Outline




- Decision Trees
- Random Forest
- Boosting
- Interpretability

# Tree-based Methods

**Objective:** Divide the predictor space into a number of simple regions


Regression Trees  
Classification Trees




CART

C4.5

Bagging & Random Forest





Boosting



# Example Tree

Predicting **the log salary** of a player as a function of the **number of hits** and **the years of experience**



# Regression Trees

$$X_1, X_2, \dots, X_p \xrightarrow{\text{distinct and nonoverlapping regions}} R_1, R_2, \dots, R_J$$

**Prediction:** Mean of the response values for the training observations in  $R_j$

$$R_1, R_2, \dots, R_J \quad ?$$

**Goal:** Finding the regions such that sum of squares is minimized

$$\sum_{j=1}^J \sum_{i \in R_j} (y_i - \hat{y}_{R_j})^2$$

$\hat{y}_{R_j}$  : mean response within  $R_j$

# Regression Trees

## Recursive Binary Splitting

$$R_1(j, s) = \{X | X_j < s\} \quad R_2(j, s) = \{X | X_j \geq s\}$$

Find  $j$  and  $s$  that minimizes

$$\sum_{i: x_i \in R_1(j, s)} (y_i - \hat{y}_{R_1})^2 + \sum_{i: x_i \in R_2(j, s)} (y_i - \hat{y}_{R_2})^2$$

### Stop

when each terminal node has a very small number of observations  
decrease in sum of squares is below a threshold (myopic)

# Regression Trees

## Tree Pruning

**Goal:** Avoiding overfitting with a fully grown or large tree  
(Selecting a subtree that leads to a lowest test error rate)

Minimize the cost  
complexity criterion: 
$$\sum_{m=1}^{|T|} \sum_{i: x_i \in R_m} (y_i - \hat{y}_{R_m})^2 + \alpha |T|$$

$T$  : subtree

$|T|$  : number of terminal nodes in  $T$


$R_m$  : region corresponding to  
the  $m$ th **terminal node**

$\alpha$  : tuning hyperparameter

$\alpha \uparrow$        $|T| \downarrow$

Use  $k$ -fold cross validation to choose  $\alpha$






# Classification Trees

Similar to a regression tree but uses different error measures based on *purity* of a region – classification is done with **majority voting**

$\hat{p}_{mk}$  : proportion of class- $k$  training observations in the  $m$ th region

## Classification Error Rate

$$E = 1 - \max_k \{\hat{p}_{mk}\}$$




## (Cross) Entropy

$$D = - \sum_{k=1}^K \hat{p}_{mk} \log \hat{p}_{mk}$$

## Gini Index

$$G = \sum_{k=1}^K \hat{p}_{mk}(1 - \hat{p}_{mk})$$




## Other Points

- **Categorical variables:** Need to consider all subsets of the possible values
  - a) Binary output – order the categories according to the proportion falling in one class, then split as if it is an ordered predictor (optimal)
  - b) Quantitative outcome – square error loss: Same as a)
  - c) Multi-category output – Trick in a) does not work (approximations)
- **Instability:** Trees have high variance due their hierarchical structure
- **CART alternatives:** ID3, C4.5, C.5.0

# Bagging

**Goal:** Use bootstrapping to grow separate *deep* trees and average all the predictions (regression) or apply majority rule (classification) to reduce the variance



## Regression

$B$  : number of separate training sets

$\hat{f}^{*b}(x)$  : prediction obtained with the  $b$ th training set

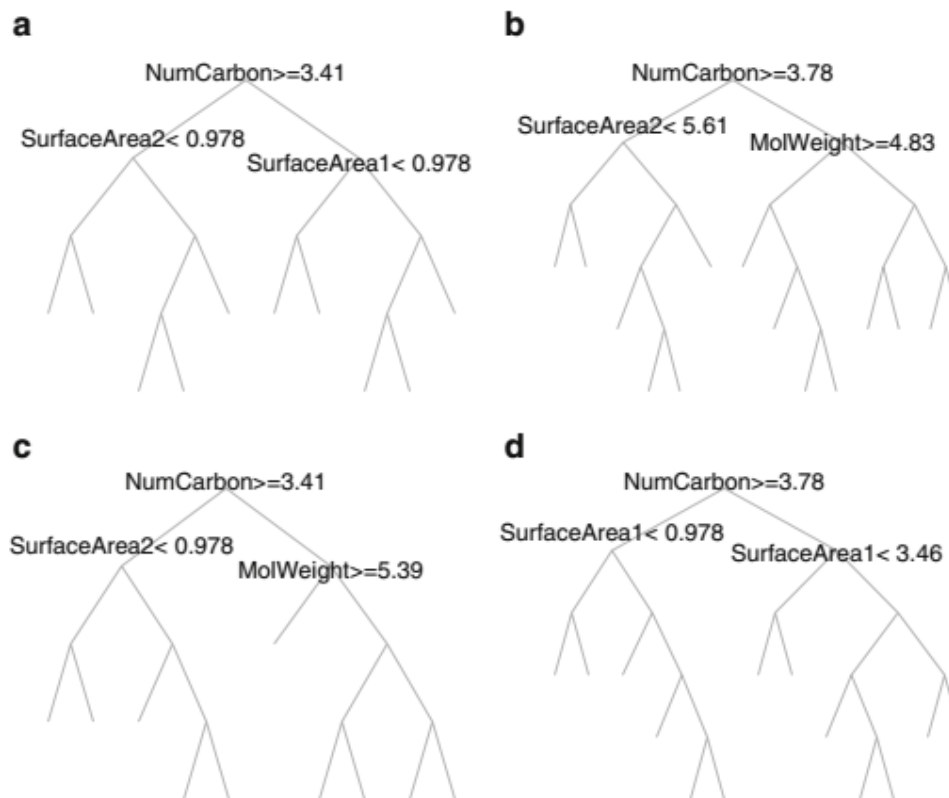
$$\hat{f}_{\text{bag}}(x) = \frac{1}{B} \sum_{b=1}^B \hat{f}^{*b}(x)$$

Classification Apply majority voting

# Random Forest

**Goal:** Smart bagging to *de-correlate* the trees – only a random sample of  $m$  predictors is chosen as candidates for splitting


Why?




\*

# Random Forest

$$m \approx \sqrt{p}$$



random subsets of size  $m$



$\{X_2, X_3, X_7\}$



## Regression

$B$  : number of grown trees


$\hat{f}^{*b}(x)$  : prediction obtained with the  $b$ th tree

$$\hat{f}_{\text{rf}}^B(x) = \sum_{b=1}^B \hat{f}^{*b}(x)$$

Classification Apply majority voting

# Boosting

**Idea:** Combining mediocre classifiers sequentially to boost their collective performance through weighted data sampling



**Update Rule:** Give more weights to **incorrectly** classified samples

## AdaBoost – Binary Classification $\{+1, -1\}$

Each sample has the same starting weight ( $1/n$ )

**for**  $k=1$  to  $K$  **do**

Fit a tree with  $d$  splits using the weighted samples and compute misclassification error ( $\epsilon_k$ )

Compute stage weight value  $\ln \frac{1 - \epsilon_k}{\epsilon_k}$

Update the sample weights – give more weights to incorrectly classified samples

**end**


Compute the boosted classifier's prediction for each sample by multiplying the  $k$ th stage value by the  $k$ th model prediction and add these quantities across  $k$ . If the sum is positive, then classify the sample as  $+1$  otherwise as  $-1$

Details **for general classifiers** are in the supplementary note on Canvas



# Notes on (Tree-based) Ensemble Methods

- Both boosting and bagging can be applied to different learning methods
- While bagging allows direct parallel implementation, the sequential structure of boosting prevents parallelization
- Ensemble methods cause loss of interpretability
- Variable importance plots can be used



**Variable importance** is computed by sorting the predictors according to the mean decrease they achieve in Gini index or entropy

(assign 100 to the largest and scale the rest)